

# IBM Data Science Capstone: Car Accident Severity Report

## 1. Introduction

Road traffic accidents are a leading cause of death in young people in the United States [1][2]. The average number of car accidents in the U.S. is 6 million car accidents every year, and about 6% of those accidents result in at least one death. 3 million people are injured as a result of car accidents and around 2 million drivers experience permanent injuries every year [3].

Analyzing historical vehicle crash data can help us understand the most common factors, including environmental conditions (weather, road surface conditions, and lighting conditions) and their correlation with accident severity. This information can be used to create a prediction model that can be used in conjunction with other Apps like Google Maps to predict the severity of an accident to help drivers be more alert to what can commonly lead to a severe accident. For this project, data from the City of Seattle's Police Department for the years 2004 until present are utilized.

## 2. Data

In this project, shared data for Seattle city from Applied Data Science Capstone Project Week1 are used [4]. The dataset consists of 38 columns, 35 columns are the attributes or independent variables. One column\* (column A and N) is the dependent or the predicted variable, SEVERITYCODE, and another column (column O) is the description of the code, SEVERITYDESC. The predicted variable has two values: either 1 for property damage only collision or 2 for injury collision. The dataset has more than 194,000 records representing all types of collisions provided by Seattle Police Department and recorded by Traffic record in the timeframe 2004 to 2020. This study aims to predict the impact of environmental conditions of the accidents, namely: WEATHER, ROADCOND, and LIGHTCOND. Brief explanation of each attribute can be found in the file uploaded to Github in the link below.

[https://github.com/Yusser89/Coursera\\_Capstone/blob/master/IBMCapstoneProjectWeek1\\_Part2.pdf](https://github.com/Yusser89/Coursera_Capstone/blob/master/IBMCapstoneProjectWeek1_Part2.pdf)

- There is a duplicate, column A and Column N both represent SEVERITYCODE

## 2.1. Feature Selection

Since the study focuses on environmental conditions of the accidents, we can narrow down the dataset to 'WEATHER', 'ROADCOND', and 'LIGHTCOND'.

We begin by importing main libraries followed by loading data file and printing the size of the dataset.

```
In [1]: #import main libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.pyplot import show
%matplotlib inline
import seaborn as sns
!pip -q install folium
import folium
```

```
In [3]: #Load data file and print size of the dataset
coll_df = pd.read_csv('https://s3.us.cloud-object-storage.appdomain.cloud/cf-
print('Dimensions of dataset:', coll_df.shape)
```

Dimensions of dataset: (194673, 38)

The dataset is comprised of 194673 records and 38 features

We can view the columns and first five rows of the dataset to get an idea of the data we are dealing with.

```
In [4]: #View column information
coll_df.columns
```

```
Out[4]: Index(['SEVERITYCODE', 'X', 'Y', 'OBJECTID', 'INCKEY', 'COLDETKEY', 'REPO
RTNO',
              'STATUS', 'ADDRTYPE', 'INTKEY', 'LOCATION', 'EXCEPTRSNCODE',
              'EXCEPTRSNDESC', 'SEVERITYCODE.1', 'SEVERITYDESC', 'COLLISIONTYP
E',
              'PERSONCOUNT', 'PEDCOUNT', 'PEDCYLCOUNT', 'VEHCOUNT', 'INCDATE',
              'INCDTTM', 'JUNCTIONTYPE', 'SDOT_COLCODE', 'SDOT_COLDESC',
              'INATTENTIONIND', 'UNDERINFL', 'WEATHER', 'ROADCOND', 'LIGHTCOND',
              'PEDROWNOTGRNT', 'SDOTCOLNUM', 'SPEEDING', 'ST_COLCODE', 'ST_COLDE
SC',
              'SEGLANEKEY', 'CROSSWALKKEY', 'HITPARKEDCAR'],
              dtype='object')
```

## By Yusser Al-Qazwini

```
In [5]: #view the first 5 rows of the dataset
coll_df.head()
```

Out[5]:

	SEVERITYCODE	X	Y	OBJECTID	INCKEY	COLDETKEY	REPORTNO
0	2	-122.323148	47.703140	1	1307	1307	3502005
1	1	-122.347294	47.647172	2	52200	52200	2607959
2	1	-122.334540	47.607871	3	26700	26700	1482393
3	1	-122.334803	47.604803	4	1144	1144	3503937
4	2	-122.306426	47.545739	5	17700	17700	1807429

5 rows × 38 columns

The target variable, 'SEVERITYCODE', is described by 'SEVERITYDESC'. Let's see how many different codes we have.

```
In [6]: #SEVERITYDESC provides a description of the severity code
print("Road severity description types : ",coll_df["SEVERITYDESC"].value_counts())

Road severity description types : Property Damage Only Collision    136485
Injury Collision                    58188
Name: SEVERITYDESC, dtype: int64
```

So we have two severity codes: 1 for property damage only collision and 2 for injury collision.

We then narrow down our dataset to the features of interest, namely: 'WEATHER', 'ROADCOND', 'LIGHTCOND'.

```
In [7]: #Since the study focuses on environmental conditions of the accidents, we can narrow down t
df = coll_df[['SEVERITYCODE', 'SEVERITYDESC', 'WEATHER', 'ROADCOND', 'LIGHTCOND']]
df.head()
```

Out[7]:

	SEVERITYCODE	SEVERITYDESC	WEATHER	ROADCOND	LIGHTCOND
0	2	Injury Collision	Overcast	Wet	Daylight
1	1	Property Damage Only Collision	Raining	Wet	Dark - Street Lights On
2	1	Property Damage Only Collision	Overcast	Dry	Daylight
3	1	Property Damage Only Collision	Clear	Dry	Daylight
4	2	Injury Collision	Raining	Wet	Daylight

```
In [8]: #print dimension of the new dataset
print('Dimensions of dataset:', df.shape)
```

Dimensions of dataset: (194673, 5)

## 2.2. Handling Missing Data

The dataset consists of raw data so there is missing information. First, we will search for question marks and replace them with NaNs. Then we will replace all NaN values with the most frequent data from each attribute. In addition to that, we are going to group some types of the features together if they are related to each other.

```
In [9]: #replace all questions marks "?" with NaN
df.replace("?", np.nan, inplace = True)
#evaluate missing data or "NaN"
missing_data = df.isnull()
for column in missing_data.columns.values.tolist():
    print(column)
    print (missing_data[column].value_counts())
    print("-----")
```

```
SEVERITYCODE
False      194673
Name: SEVERITYCODE, dtype: int64
-----
SEVERITYDESC
False      194673
Name: SEVERITYDESC, dtype: int64
-----
WEATHER
False      189592
True        5081
Name: WEATHER, dtype: int64
-----
ROADCOND
False      189661
True        5012
Name: ROADCOND, dtype: int64
-----
LIGHTCOND
False      189503
True        5170
Name: LIGHTCOND, dtype: int64
-----
```

From the results above, it can be seen that we are missing 5081 weather data, 5012 road condition data, and 5170 light condition data. This missing information needs to be addressed.

Let's also explore the different types of each feature to see if we can group them together

```
In [10]: #view weather conditions
print("Weather conditions reported are: ")
print(df['WEATHER'].value_counts())
print("-----")
#view road conditions
print("Road conditions reported are: ")
print(df['ROADCOND'].value_counts())
print("-----")
#view light conditions
print("Road conditions reported are: ")
print(df['LIGHTCOND'].value_counts())
print("-----")
```

```
Weather conditions reported are:
Clear                111135
Raining              33145
Overcast             27714
Unknown              15091
Snowing              907
Other                 832
Fog/Smog/Smoke       569
Sleet/Hail/Freezing Rain 113
Blowing Sand/Dirt     56
Severe Crosswind      25
Partly Cloudy         5
Name: WEATHER, dtype: int64
```

```
-----
Road conditions reported are:
Dry                  124510
Wet                  47474
Unknown              15078
Ice                  1209
Snow/Slush           1004
Other                132
Standing Water       115
Sand/Mud/Dirt        75
Oil                  64
Name: ROADCOND, dtype: int64
```

```
-----
Light conditions reported are:
Daylight             116137
Dark - Street Lights On 48507
Unknown              13473
Dusk                  5902
Dawn                  2502
Dark - No Street Lights 1537
Dark - Street Lights Off 1199
Other                 235
Dark - Unknown Lighting 11
Name: LIGHTCOND, dtype: int64
-----
```

Weather conditions can be grouped as follows:

- SevereWeather: Raining, Snowing, Sleet/Hail/Freezing Rain, Fog/Smog/Smoke,
- Blowing Sand/Dirt, Severe Crosswind
- Overcast: PartlyCloudy and Overcast
- Unknown: Other

```
In [11]: #replacing missing values with the most frequent ones and grouping the similar types together
df["WEATHER"].replace(np.nan, df["WEATHER"].value_counts().idxmax(), inplace=True)
df["WEATHER"].replace("Other", "Unknown", inplace=True)
df["WEATHER"].replace("Partly Cloudy", "Overcast", inplace=True)
df["WEATHER"].replace("Raining", "SevereWeather", inplace=True)
df["WEATHER"].replace("Snowing", "SevereWeather", inplace=True)
df["WEATHER"].replace("Sleet/Hail/Freezing Rain", "SevereWeather", inplace=True)
df["WEATHER"].replace("Fog/Smog/Smoke", "SevereWeather", inplace=True)
df["WEATHER"].replace("Blowing Sand/Dirt", "SevereWeather", inplace=True)
df["WEATHER"].replace("Severe Crosswind", "SevereWeather", inplace=True)
df["WEATHER"].value_counts()
```

```
Out[11]: Clear          116216
SevereWeather         34815
Overcast              27719
Unknown               15923
Name: WEATHER, dtype: int64
```

Road conditions can be grouped as follows:

- IceOilWaterSnow: Ice, Standing Water, Oil, Snow/Slush, Sand/Mud/Dirt
- Unknown: Other

```
In [12]: #replacing missing values with the most frequent ones and grouping the similar types together
df["ROADCOND"].replace(np.nan, df["ROADCOND"].value_counts().idxmax(), inplace=True)
df["ROADCOND"].replace("Ice", "IceOilWaterSnow", inplace=True)
df["ROADCOND"].replace("Standing Water", "IceOilWaterSnow", inplace=True)
df["ROADCOND"].replace("Oil", "IceOilWaterSnow", inplace=True)
df["ROADCOND"].replace("Snow/Slush", "IceOilWaterSnow", inplace=True)
df["ROADCOND"].replace("Other", "Unknown", inplace=True)
df["ROADCOND"].replace("Sand/Mud/Dirt", "IceOilWaterSnow", inplace=True)
df["ROADCOND"].value_counts()
```

```
Out[12]: Dry           129522
Wet                 47474
Unknown             15210
IceOilWaterSnow      2467
Name: ROADCOND, dtype: int64
```

Light conditions can be grouped as follows:

- Dark-No-Light: Dark — No Street Lights, Dark — Street Lights Off, Dark — Unknown Lighting
- Dark-With-Light: Dark — Street Lights On  
DuskDawn: Dusk, Dawn
- Unknown: Other

```
In [13]: #replacing missing values with the most frequent ones and grouping the similar types together
df["LIGHTCOND"].replace(np.nan, df['LIGHTCOND'].value_counts().idxmax(), inplace=True)
df["LIGHTCOND"].replace("Dark - No Street Lights", "Dark-No-Light", inplace=True)
df["LIGHTCOND"].replace("Dark - Street Lights Off", "Dark-No-Light", inplace=True)
df["LIGHTCOND"].replace("Dark - Unknown Lighting", "Dark-No-Light", inplace=True)
df["LIGHTCOND"].replace("Dark - Street Lights On", "Dark-With-Light", inplace=True)
df["LIGHTCOND"].replace("Other", "Unknown", inplace=True)
df["LIGHTCOND"].replace("Dusk", "DuskDawn", inplace=True)
df["LIGHTCOND"].replace("Dawn", "DuskDawn", inplace=True)
df['LIGHTCOND'].value_counts()

Out[13]: Daylight          121307
Dark-With-Light         48507
Unknown                 13708
DuskDawn                 8404
Dark-No-Light           2747
Name: LIGHTCOND, dtype: int64
```

Let's check if we have any null values

```
In [14]: #display dataset information: attribute name, count, how many are null, and data type
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 194673 entries, 0 to 194672
Data columns (total 5 columns):
SEVERITYCODE      194673 non-null int64
SEVERITYDESC      194673 non-null object
WEATHER           194673 non-null object
ROADCOND          194673 non-null object
LIGHTCOND         194673 non-null object
dtypes: int64(1), object(4)
memory usage: 7.4+ MB
```

### 3. Methodology

In this section of the report, exploratory data analysis, inferential statistical testing, and machine learnings used are described.

#### 3.1. Data Visualization


Number of accidents are plotted against each environmental factor (feature) with percentage of each type of each feature to understand the impact of each factor.

First let's see the impact of **weather conditions**.

```
In [15]: #Weather conditions
plt.figure(figsize=(15,8))
sns.set(style="white")
custom_palette = ["red", "grey"]
sns.set_palette(custom_palette)
total = float(len(df)) # one person per row
ax = sns.countplot(x="WEATHER", hue="SEVERITYDESC", data=df) # for Seaborn version 0.7 and
ax.set(xlabel="Weather Condition", ylabel = "Number of Accidents")

for p in ax.patches:
    height = p.get_height()
    ax.text(p.get_x()+p.get_width()/2.,
            height + 3,
            '{:1.2%}'.format(height/total),
            ha="center")

show()
```




We can see from the graph above that majority of the accidents happened in clear weather. I was expecting to see more accidents in severe weather. We need more information on 'Unknown' weather conditions as the percentage should not be neglected particularly for accidents that caused property damage only.

Let's now see the impact of **road conditions**.



## By Yusser Al-Qazwini


```
In [16]: #road conditions
plt.figure(figsize=(15,8))
sns.set(style="white")
custom_palette = ["red","grey"]
sns.set_palette(custom_palette)
total = float(len(df)) # one person per row
#ax = sns.barplot(x="class", hue="who", data=titanic)
ax = sns.countplot(x="ROADCOND", hue="SEVERITYDESC", data=df) # for Seaborn version 0.7 and
ax.set(xlabel="Road Condition", ylabel = "Number of Accidents")
for p in ax.patches:
    height = p.get_height()
    ax.text(p.get_x()+p.get_width()/2.,
            height + 3,
            '{:1.2%}'.format(height/total),
            ha="center")
show()
```



We can see from the graph above that majority of the accidents happened on dry roads. I was expecting to see more accidents on wet or icy, snowy, oily roads! We also need more information on 'Unknown' road conditions as the percentage should not be neglected particularly for accidents that caused property damage only.

And finally let's examine the impact of **light conditions**.

```
In [17]: #light conditions
plt.figure(figsize=(15,8))
sns.set(style="white")
custom_palette = ["red","grey"]
sns.set_palette(custom_palette)
total = float(len(df)) # one person per row
#ax = sns.barplot(x="class", hue="who", data=titanic)
ax = sns.countplot(x="LIGHTCOND", hue="SEVERITYDESC", data=df) # for Seaborn version 0.7 and above
ax.set(xlabel="Light Condition", ylabel = "Number of Accidents")
for p in ax.patches:
    height = p.get_height()
    ax.text(p.get_x()+p.get_width()/2.,
            height + 3,
            '{:1.2%}'.format(height/total),
            ha="center")
show()
```



It can be seen from the graph above that majority of accidents happened during the day with daylight. This also was not as I expected! Again, we need more information on 'Unknown' light conditions as the percentage should not be neglected particularly for accidents that caused property damage only.

## 3.2 Machine Learning Model Selection

The preprocessed dataset can be split into training and test sub datasets (70% for training and 30% for testing) using the scikit learn "train\_test\_split" method. Since the target column (SEVERITYCODE) is categorical, a classification model is used to predict the severity of an accident. Three classification models were trained and evaluated, namely: K-Nearest Neighbor, Decision Tree, and Logistic Regression.

We will start by defining the X (independent variables) and y (dependent variable) as follows.

```
In [18]: #SEVERITYCODE is the dependent variable to be predicted and it is already numerical  
y = np.asarray(df["SEVERITYCODE"])  
y[0:5]
```

```
Out[18]: array([2, 1, 1, 1, 2])
```

```
In [19]: #defining X as the enviromental features  
Feature = df[['WEATHER', 'ROADCOND', 'LIGHTCOND']]  
X = Feature.values  
X[0:5]
```

```
Out[19]: array(['Overcast', 'Wet', 'Daylight',  
               ['SevereWeather', 'Wet', 'Dark-With-Light'],  
               ['Overcast', 'Dry', 'Daylight'],  
               ['Clear', 'Dry', 'Daylight'],  
               ['SevereWeather', 'Wet', 'Daylight']], dtype=object)
```

X data needs to be converted to numerical data to be used in the classification models. This can be achieved by using Label Encoding.

```
In [20]: #categorical boolean mask  
categorical_feature_mask = Feature.dtypes==object  
categorical_feature_mask
```

```
Out[20]: WEATHER      True  
         ROADCOND    True  
         LIGHTCOND   True  
         dtype: bool
```

```
In [21]: #filter categorical columns using mask and turn it into a List  
categorical_cols = Feature.columns[categorical_feature_mask].tolist()  
categorical_cols
```

```
Out[21]: ['WEATHER', 'ROADCOND', 'LIGHTCOND']
```

```
In [22]: #import LabelEncoder to convert each class under specified feature to a numerical value.  
from sklearn.preprocessing import LabelEncoder  
# instantiate LabelEncoder object  
le = LabelEncoder()
```

```
In [23]: #apply le on categorical feature columns  
Feature[categorical_cols] = Feature[categorical_cols].apply(lambda col: le.fit_transform(col))  
Feature[categorical_cols].head(10)
```

Out[23]:

	WEATHER	ROADCOND	LIGHTCOND
0	1	3	2
1	2	3	1
2	1	0	2
3	0	0	2
4	2	3	2
5	0	0	2
6	2	3	2
7	0	0	2
8	0	0	2
9	0	0	2

```
In [24]: #display the feature list as numerical values
X = Feature.values
X[0:5]
```

```
Out[24]: array([[1, 3, 2],
                [2, 3, 1],
                [1, 0, 2],
                [0, 0, 2],
                [2, 3, 2]])
```

It is always better to normalize the features data.

```
In [25]: #normalize the feature data
from sklearn import preprocessing
X = preprocessing.StandardScaler().fit(X).transform(X)
X[0:5]
```

```
Out[25]: array([[ 0.24982869,  1.6059602,  0.11844397],
                [ 1.23126144,  1.6059602, -1.14979445],
                [ 0.24982869, -0.68884852,  0.11844397],
                [-0.73160406, -0.68884852,  0.11844397],
                [ 1.23126144,  1.6059602,  0.11844397]])
```

### 3.2.1. Model

It's time to build our models by first splitting our data into training and testing sets of 70% and 30% respectively.

```
In [60]: #Test/Train split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=4)
print ('Train set:', X_train.shape, y_train.shape)
print ('Test set:', X_test.shape, y_test.shape)
```

```
Train set: (136271, 3) (136271,)
Test set: (58402, 3) (58402,)
```

### 3.2.1.1. K Nearest Neighbor (KNN)

KNN is used to predict the severity of an accident of an unknown dataset based on its proximity in the multi-dimensional hyperspace of the feature set to its “k” nearest neighbors, which have known outcomes. Since finding the best k is memory-consuming and time-consuming, we will use k=25 based on [5].

```
In [61]: from sklearn.neighbors import KNeighborsClassifier
k = 25
#Train Model and Predict
neigh = KNeighborsClassifier(n_neighbors = k).fit(X_train,y_train)
neigh

Out[61]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=None, n_neighbors=25, p=2,
                             weights='uniform')
```

```
In [63]: #train dataset
Kyhat_train = neigh.predict(X_train)
Kyhat_train[0:5]

Out[63]: array([1, 1, 1, 1, 1])
```

```
In [86]: #test dataset
Kyhat_test = neigh.predict(X_test)
Kyhat_test[0:5]

Out[86]: array([1, 1, 1, 1, 1])
```

### 3.2.1.2. Decision Tree

A decision tree model is built from historical data of accident severity in relationship to environmental conditions. Then the trained decision tree can be used to predict the severity of an accident. Since finding the maximum depth is also memory and time consuming, will use max\_depth=30 based on [5].

```
In [66]: from sklearn.tree import DecisionTreeClassifier
SevTree = DecisionTreeClassifier(criterion="entropy", max_depth = 30)
SevTree.fit(X_train,y_train)
SevTree # it shows the default parameters

Out[66]: DecisionTreeClassifier(class_weight=None, criterion='entropy', max_depth=30,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                                splitter='best')
```

```
In [69]: #train dataset
dyhat_train = SevTree.predict(X_train)
dyhat_train[0:5]
```

```
Out[69]: array([1, 1, 1, 1, 1])
```

```
In [84]: #test dataset
dyhat_test = SevTree.predict(X_test)
dyhat_test[0:5]
```

```
Out[84]: array([1, 1, 1, 1, 1])
```

### 3.2.1.3. Logistic Regression

Logistic Regression is useful when the observed dependent variable,  $y$ , is categorical. It produces a formula that predicts the probability of the class label as a function of the independent variables. An inverse-regularisation strength of  $C=0.01$  is used as in [5].

```
In [71]: from sklearn.linear_model import LogisticRegression
LR = LogisticRegression(C=0.01, solver='liblinear').fit(X_train,y_train)
LR
```

```
Out[71]: LogisticRegression(C=0.01, class_weight=None, dual=False, fit_intercept=True,
                           intercept_scaling=1, max_iter=100, multi_class='warn',
                           n_jobs=None, penalty='l2', random_state=None, solver='liblinear',
                           tol=0.0001, verbose=0, warm_start=False)
```

```
In [80]: #train dataset
lyhat_train = LR.predict(X_train)
lyhat_train[0:5]
```

```
Out[80]: array([1, 1, 1, 1, 1])
```

```
In [87]: #test dataset
lyhat_test = LR.predict(X_test)
lyhat_test[0:5]
```

```
Out[87]: array([1, 1, 1, 1, 1])
```

## 4. Results (Model Evaluation)

Accuracy of the 3 models is calculated using these metrics: Jaccard Similarity Score, F1-SCORE, and LOGLOSS (with Linear Regression).

```
In [110]: from sklearn import metrics
from sklearn.metrics import log_loss
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import jaccard_similarity_score

print('Jaccard Similarity Score: ')
print('KNN Model: ', jaccard_similarity_score(y_test,Kyhat_test))
print('Decision Tree: ', jaccard_similarity_score(y_test,dyhat_test))
print('Logistic Regression: ', jaccard_similarity_score(y_test,lyhat_test))
print('-----')
print('F1-SCORE: ')
print('KNN Model: ', metrics.f1_score(y_test,Kyhat_test, average='weighted'))
print('Decision Tree: ', metrics.f1_score(y_test,dyhat_test, average='weighted'))
print('Logistic Regression: ', metrics.f1_score(y_test,lyhat_test, average='weighted'))
print('-----')
print('LOGLOSS for Logistic Regression: ')
lyhat_test_prob=LR.predict_proba(X_test)
print(log_loss(y_test,lyhat_test_prob))
```

```
Jaccard Similarity Score:
KNN Model:  0.6941543097839115
Decision Tree:  0.7034348138762371
Logistic Regression:  0.7034519365775145
-----
F1-SCORE:
KNN Model:  0.5912085352895935
Decision Tree:  0.5809821168927154
Logistic Regression:  0.5809904188654375
-----
LOGLOSS for Logistic Regression:
0.5991064490814039
```

## 5. Discussion

First the dataset had categorical data of type 'object'. Label encoding was used to convert categorical features to numerical values. The imbalanced data issue was ignored because there was a problem installing imbalanced-learn to use imblearn.

Once data was cleaned and analyzed, it was fed into three ML models: K-Nearest Neighbor, Decision Tree, and Logistic Regression. Values of k, max depth and inverse-regularisation strength C were taken from [5]. Evaluation metrics used to test the accuracy of the models were Jaccard Similarity Index, F-1 SCORE and LOGLOSS for Logistic Regression.

It is highly recommended to solve the data imbalance problem for more accurate results.

## 6. Conclusion

The goal of this project is to analyze historical vehicle crash data to understand the correlation of environmental conditions (weather, road surface, and lighting conditions) with accident severity. Vehicle accident data from the City of Seattle's Police Department

for the years 2004 until present were Used. The data was cleaned, and features related to environmental conditions were selected and analyzed. It was found that majority of accidents happened in clear weather, dry roads, and during daytime which wasn't what I expected. Machine learning models; K-Nearest Neighbor, Decision Tree and Logistic Regression were used to predict the severity of an accident based on certain environmental conditions. The models used were also evaluated using different accuracy metrics.

## **7. References**

1. Road Traffic Injuries and Deaths — A Global Problem. CDC, Center for Disease Control and Prevention, <https://www.cdc.gov/injury/features/global-road-safety/index.html#:~:text=Road%20traffic%20crashes%20are%20a,citizens%20residing%20or%20traveling%20abroad.>
2. Road Traffic Injuries. WHO, Global Health Observation Data, [https://www.who.int/health-topics/road-safety#tab=tab\\_1](https://www.who.int/health-topics/road-safety#tab=tab_1)
3. Car Accident Statistics in the U.S. Driver Knowledge, <https://www.driverknowledge.com/car-accident-statistics/#:~:text=U.S.%20every%20year%20is%206,experience%20permanent%20injuries%20every%20year>
4. Shared data for Seattle city from Applied Data Science Capstone Project Week1, <https://s3.us.cloud-object-storage.appdomain.cloud/cf-courses-data/CognitiveClass/DP0701EN/version-2/Data-Collisions.csv>
5. Seattle Car Accident Severity — IBM Capstone Project by AP Thomson, <https://medium.com/@alasdair.p.thomson/seattle-car-accident-severity-ibm-capstone-project-9cef20fc7e6adn>